

2.0-raa-train-model

May 31, 2025

```
[1]: %load_ext autoreload
      %autoreload 2
```

```
[ ]: import os
      import tensorflow as tf
```

```
[4]: dataset_dir = os.path.join('data', 'dataloader')
```

```
[ ]: img_height = 224
      img_width = 224
      batch_size = 32
```

```
[ ]: train_ds = tf.keras.preprocessing.image_dataset_from_directory(
      dataset_dir,
      validation_split=0.2,
      subset="training",
      seed=42,
      image_size=(img_height, img_width),
      batch_size=batch_size
  )

  val_ds = tf.keras.preprocessing.image_dataset_from_directory(
      dataset_dir,
      validation_split=0.2,
      subset="validation",
      seed=42,
      image_size=(img_height, img_width),
      batch_size=batch_size
  )
```


```
Found 19833 files belonging to 12 classes.
Using 15867 files for training.
Found 19833 files belonging to 12 classes.
Using 3966 files for validation.
```

```
[7]: class_names = train_ds.class_names
      print(f"Clases detectadas: {class_names}")
```

Clases detectadas: ['bird', 'bobcat', 'car', 'cat', 'coyote', 'dog', 'opossum', 'rabbit', 'raccoon', 'rodent', 'skunk', 'squirrel']

```
[8]: import matplotlib.pyplot as plt
```

```
[9]: plt.figure(figsize=(10, 10))  
for images, labels in train_ds.take(1):  
    for i in range(9):  
        ax = plt.subplot(3, 3, i + 1)  
        plt.imshow(images[i].numpy().astype("uint8"))  
        plt.title(class_names[labels[i]])  
        plt.axis("off")
```



0.1 Justificación del enfoque experimental

Aunque hemos preparado un dataset con casi 20k recortes, sigue siendo **relativamente pequeño** para entrenar redes neuronales profundas (DNNs) desde cero. Sería idóneo utilizar otras arquitecturas.

¿Por qué?

- Modelos como ResNet, EfficientNet o Vision Transformers se entrenan sobre datasets enormes como **ImageNet** (1M+ imágenes).
- Las DNNs requieren grandes volúmenes de datos para evitar underfitting y aprender desde patrones básicos (bordes) hasta complejos (formas).
- En datasets pequeños, la práctica común es usar **transfer learning**, aprovechando pesos preentrenados y ajustando solo las capas finales.

0.1.1 Referencias relevantes

- [EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks \(arXiv\)](#)
- [Deep Residual Learning for Image Recognition — ResNet \(arXiv\)](#)
- [Keras Applications: ResNet Models \(Keras.io\)](#)
- [TensorFlow Tutorial: Transfer Learning \(TensorFlow.org\)](#)
- [CS231n Stanford Lecture 17: Transfer Learning and Domain Adaptation \(PDF\)](#)

Sin embargo, antes de saltar al transfer learning, comprobaremos **empíricamente** qué rendimiento podemos lograr construyendo una CNN sencilla desde cero.

Esto nos dará:

- Una línea base de rendimiento.

- Comprensión de las limitaciones del entrenamiento puro.

- Datos para justificar (o no) el uso de modelos preentrenados.

Próximo paso:

Construiremos y entrenaremos una CNN simple sobre nuestro dataset y analizaremos los resultados antes de decidir si migrar a transferencia.

```
[10]: from tensorflow.keras import layers, models
```

```
[12]: img_height = 224
      img_width = 224
      batch_size = 32
      epochs = 20
      num_classes = len(class_names)
```

```
[16]: earlystop_callback = tf.keras.callbacks.EarlyStopping(
      monitor='val_loss',
```

```

    patience=3,
    restore_best_weights=True
)

```

```

[20]: model = models.Sequential([
        layers.Rescaling(1./255, input_shape=(img_height, img_width, 3)),

        layers.Conv2D(32, 3, activation='relu'),
        layers.MaxPooling2D(),
        layers.BatchNormalization(),

        layers.Conv2D(64, 3, activation='relu'),
        layers.MaxPooling2D(),
        layers.BatchNormalization(),

        layers.Conv2D(128, 3, activation='relu'),
        layers.MaxPooling2D(),
        layers.BatchNormalization(),

        layers.Flatten(),
        layers.Dense(128, activation='relu'),
        layers.Dropout(0.5),
        layers.Dense(num_classes, activation='softmax')
    ])

```

```

[21]: model.compile(
        optimizer='adam',
        loss='sparse_categorical_crossentropy',
        metrics=['accuracy']
    )

```

```

[22]: model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
rescaling_3 (Rescaling)	(None, 224, 224, 3)	0
conv2d_6 (Conv2D)	(None, 222, 222, 32)	896
max_pooling2d_4 (MaxPooling2D)	(None, 111, 111, 32)	0
batch_normalization_4 (BatchNormalization)	(None, 111, 111, 32)	128

conv2d_7 (Conv2D)	(None, 109, 109, 64)	18,496
max_pooling2d_5 (MaxPooling2D)	(None, 54, 54, 64)	0
batch_normalization_5 (BatchNormalization)	(None, 54, 54, 64)	256
conv2d_8 (Conv2D)	(None, 52, 52, 128)	73,856
max_pooling2d_6 (MaxPooling2D)	(None, 26, 26, 128)	0
batch_normalization_6 (BatchNormalization)	(None, 26, 26, 128)	512
flatten_1 (Flatten)	(None, 86528)	0
dense_1 (Dense)	(None, 128)	11,075,712
dropout (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 12)	1,548

Total params: 11,171,404 (42.62 MB)

Trainable params: 11,170,956 (42.61 MB)

Non-trainable params: 448 (1.75 KB)

```
[23]: history = model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=epochs,
    callbacks=[earlystop_callback]
)
```

Epoch 1/20

496/496 403s 803ms/step -

accuracy: 0.2463 - loss: 20.2863 - val_accuracy: 0.2239 - val_loss: 2.4814

Epoch 2/20

496/496 378s 761ms/step -

accuracy: 0.2228 - loss: 3.0947 - val_accuracy: 0.2665 - val_loss: 2.1480

Epoch 3/20

496/496 387s 780ms/step -

accuracy: 0.2435 - loss: 2.3733 - val_accuracy: 0.2506 - val_loss: 4.1784

```

Epoch 4/20
496/496          385s 775ms/step -
accuracy: 0.2533 - loss: 2.3134 - val_accuracy: 0.2774 - val_loss: 2.9400
Epoch 5/20
496/496          381s 768ms/step -
accuracy: 0.2584 - loss: 2.1813 - val_accuracy: 0.2708 - val_loss: 2.1100
Epoch 6/20
496/496          385s 777ms/step -
accuracy: 0.2652 - loss: 2.1524 - val_accuracy: 0.2877 - val_loss: 2.1066
Epoch 7/20
496/496          392s 790ms/step -
accuracy: 0.2598 - loss: 2.1828 - val_accuracy: 0.2988 - val_loss: 2.1902
Epoch 8/20
496/496          373s 751ms/step -
accuracy: 0.2639 - loss: 2.2735 - val_accuracy: 0.2680 - val_loss: 2.0997
Epoch 9/20
496/496          335s 676ms/step -
accuracy: 0.2739 - loss: 2.0743 - val_accuracy: 0.3051 - val_loss: 2.1709
Epoch 10/20
496/496          352s 710ms/step -
accuracy: 0.2744 - loss: 2.0924 - val_accuracy: 0.1639 - val_loss: 6.6888
Epoch 11/20
496/496          359s 724ms/step -
accuracy: 0.2676 - loss: 2.1081 - val_accuracy: 0.3316 - val_loss: 2.0265
Epoch 12/20
496/496          355s 715ms/step -
accuracy: 0.2732 - loss: 2.0800 - val_accuracy: 0.3180 - val_loss: 1.9643
Epoch 13/20
496/496          357s 720ms/step -
accuracy: 0.2797 - loss: 2.0451 - val_accuracy: 0.2872 - val_loss: 2.4930
Epoch 14/20
496/496          356s 718ms/step -
accuracy: 0.2888 - loss: 2.0273 - val_accuracy: 0.2753 - val_loss: 2.2652
Epoch 15/20
496/496          373s 751ms/step -
accuracy: 0.2807 - loss: 2.0677 - val_accuracy: 0.2388 - val_loss: 2.3134

```

```

[5]: def plot_training_history(history, title_prefix=""):
    """
    Grafica loss y accuracy de entrenamiento y validación.

    Args:
        history: history object (keras.callbacks.History) después de model.
        ↪fit().
        title_prefix: (opcional) texto para añadir al inicio de los títulos de
        ↪las gráficas.
    """

```

```

plt.figure(figsize=(12, 5))


# Loss
plt.subplot(1, 2, 1)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss',
color='orange')
plt.title(f'{title_prefix} Loss: Train vs Validation')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

# Accuracy
plt.subplot(1, 2, 2)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy',
color='orange')
plt.title(f'{title_prefix} Accuracy: Train vs Validation')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout()
plt.show()

```

[34]: `plot_training_history(history, title_prefix="Model 1: ")`



Diagnóstico

- El modelo ya tocó su límite de aprendizaje.
- No hay mejora significativa con más epochs.
- Hay alta varianza en validación; podría ser underfitting, mala arquitectura, o ambos.

Seguir con esta arquitectura, solo si disponemos de más datos

0.1.2 Aumentando dataset

```
[16]: from vision.config import DATALOADER, ANNOTATIONS
```

```
[17]: import json
import pandas as pd
import os
```

```
[18]: json_files = [
    'cis_test_annotations.json',
    'cis_val_annotations.json',
    'trans_test_annotations.json',
    'trans_val_annotations.json',
    'train_annotations.json'
]
```

```
[19]: from tqdm import tqdm
```

```
[20]: def load_and_merge_annotations(base_path, json_files):
    merged_dfs = []

    for file in tqdm(json_files, desc="Processing JSON files"):
        file_path = os.path.join(base_path, file)
        with open(file_path, 'r') as f:
            data = json.load(f)

        images_df = pd.DataFrame(data['images'])
        categories_df = pd.DataFrame(data['categories'])
        annotations_df = pd.DataFrame(data['annotations'])

        merged_df = annotations_df.merge(
            images_df, left_on='image_id', right_on='id', suffixes=('_ann',
↵ '_img')
        )
        merged_df = merged_df.merge(
            categories_df, left_on='category_id', right_on='id', suffixes=('',
↵ '_cat')
        )

        merged_dfs.append(merged_df)

    final_df = pd.concat(merged_dfs, ignore_index=True)
    return final_df
```



```
[21]: combined_df = load_and_merge_annotations(ANNOTATIONS, json_files)
```

```
Processing JSON files: 100%|          | 5/5 [00:00<00:00,  7.77it/s]
```

```
[22]: combined_df.head()
```

```
[22]:
```

	image_id	category_id	\
0	5a263af8-23d2-11e8-a6a3-ec086b02610b	30	
1	5971faf4-23d2-11e8-a6a3-ec086b02610b	30	
2	59ffbd00-23d2-11e8-a6a3-ec086b02610b	1	
3	599d7b4a-23d2-11e8-a6a3-ec086b02610b	1	
4	599fbfe7-23d2-11e8-a6a3-ec086b02610b	30	

	id_ann	\
0	343db02b-7d5b-11e7-884d-7845c41c2c67	
1	343db02c-7d5b-11e7-884d-7845c41c2c67	
2	36132	
3	11320	
4	343db02f-7d5b-11e7-884d-7845c41c2c67	

	bbox	\
0	NaN	
1	NaN	
2	[1118.72, 570.88, 328.96000000000004, 180.4800...	
3	[513.28, 500.47999999999996, 153.60000000000000...	
4	NaN	

	file_name	rights_holder	height	width	\
0	5a263af8-23d2-11e8-a6a3-ec086b02610b.jpg	Justin Brown	1494	2048	
1	5971faf4-23d2-11e8-a6a3-ec086b02610b.jpg	Justin Brown	1494	2048	
2	59ffbd00-23d2-11e8-a6a3-ec086b02610b.jpg	Justin Brown	1494	2048	
3	599d7b4a-23d2-11e8-a6a3-ec086b02610b.jpg	Justin Brown	1494	2048	
4	599fbfe7-23d2-11e8-a6a3-ec086b02610b.jpg	Justin Brown	1494	2048	

	frame_num	date_captured	location	seq_num_frames	\
0	2	2012-01-29 12:35:56	38	3	
1	1	2012-01-23 16:56:32	38	3	
2	1	2012-02-11 01:48:19	38	3	
3	3	2012-02-01 12:24:08	38	3	
4	2	2012-02-08 03:21:39	38	3	

	seq_id	id_img	\
0	6f026d3d-5567-11e8-960a-dca9047ef277	5a263af8-23d2-11e8-a6a3-ec086b02610b	
1	6f025d61-5567-11e8-85d1-dca9047ef277	5971faf4-23d2-11e8-a6a3-ec086b02610b	
2	6f0298b8-5567-11e8-8804-dca9047ef277	59ffbd00-23d2-11e8-a6a3-ec086b02610b	
3	6f0278ba-5567-11e8-8041-dca9047ef277	599d7b4a-23d2-11e8-a6a3-ec086b02610b	
4	6f028e68-5567-11e8-a478-dca9047ef277	599fbfe7-23d2-11e8-a6a3-ec086b02610b	

	id	name
0	30	empty
1	30	empty
2	1	opossum
3	1	opossum
4	30	empty

```
[23]: from vision.dataset import rescale_bounding_boxes
```

```
[24]: rescaled_df = rescale_bounding_boxes(combined_df, target_width=1024,
      ↪target_height=747)
rescaled_df[['file_name', 'bbox_scaled']].head()
```

```
[24]:
```

	file_name \	bbox_scaled
0	5a263af8-23d2-11e8-a6a3-ec086b02610b.jpg	None
1	5971faf4-23d2-11e8-a6a3-ec086b02610b.jpg	None
2	59ffbd00-23d2-11e8-a6a3-ec086b02610b.jpg	[559.36, 285.44, 164.48000000000002, 90.240000...
3	599d7b4a-23d2-11e8-a6a3-ec086b02610b.jpg	[256.64, 250.23999999999998, 76.80000000000001...
4	599fbfe7-23d2-11e8-a6a3-ec086b02610b.jpg	None

```
[25]: from vision.plots import show_random_image_with_bbox, show_image_with_multi_bbox
      from vision.config import RAW
```

```
[45]: show_random_image_with_bbox(rescaled_df, RAW)
```

Category: coyote



```
[48]: show_image_with_multi_bbox(rescaled_df, RAW)
```



```
[49]: valid_df = rescaled_df[rescaled_df['bbox_scaled'].notnull()]
      class_counts = valid_df['name'].value_counts()
      class_counts
```

```
[49]: name
      opossum      11970
      raccoon       7076
      rabbit       4995
      coyote       4700
      bobcat       4320
      cat         4303
      dog         2641
      car         2613
      squirrel    2479
      bird        1579
      skunk        722
      rodent       482
      deer        186
      badger       29
      fox          8
```

Name: count, dtype: int64

0.1.3 Análisis de balance de clases

Además de excluir fox, badger y deer por bajo número de ejemplos, será bueno también excluir rodent (218) y skunk (508), porque siguen siendo muy minoritarios frente a las clases principales.

También se aplicara *data augmentation*

```
[50]: from vision.dataset import save_recortes_by_class
      from vision.config import RAW, CLIPS
```

```
[51]: save_recortes_by_class(valid_df, RAW, CLIPS)
```

Saving crops: 100%| | 48103/48103 [20:41<00:00, 38.73it/s]

```
[52]: from vision.dataset import count_files_per_class
```

```
[53]: counts = count_files_per_class(CLIPS)
      counts
```

```
[53]: {'opossum': 11970,
      'raccoon': 7076,
      'rabbit': 4995,
      'coyote': 4700,
      'bobcat': 4320,
      'cat': 4303,
      'dog': 2641,
      'car': 2613,
      'squirrel': 2479,
      'bird': 1579,
      'skunk': 722,
      'rodent': 482,
      'deer': 186,
      'badger': 29,
      'fox': 8}
```

Clases a excluir: fox, badger, deer, rodent, skunk

además...

```
[54]: from vision.plots import show_random_image_by_class
```

```
[55]: show_random_image_by_class(valid_df, 'car', RAW)
```

Category: car



En datasets de fauna (como cámaras trampa) es común encontrar categorías “no-animal” como car, human, vegetation, etc.

¿Por qué? Porque esos elementos aparecen accidentalmente en el entorno.

Conviene mantener clases distractoras (como car, human, etc.) porque enseñan al modelo a diferenciar fauna de objetos no relevantes, evitando falsos positivos y mejorando la generalización.

Aplicando transformaciones

```
[61]: import albumentations as A
```

```
[62]: augment_dog_car_squirrel = A.Compose([
    A.HorizontalFlip(p=0.5),
    A.RandomBrightnessContrast(p=0.3),
    A.Affine(translate_percent=0.05, scale=1.1, rotate=10, p=0.4),
    A.GaussNoise(p=0.3),
])
```

```
[63]: augment_bird = A.Compose([
        A.HorizontalFlip(p=0.5),
        A.RandomBrightnessContrast(p=0.4),
        A.Rotate(limit=20, p=0.6),
        A.Affine(translate_percent=0.1, scale=1.2, rotate=15, p=0.5),
        A.GaussNoise(p=0.4),
        A.Blur(blur_limit=3, p=0.2),
    ])
```

```
[64]: from vision.dataset import apply_augmentations_for_class
```

```
[65]: from vision.config import AUGMENTED
```

```
[66]: apply_augmentations_for_class('dog', CLIPS, AUGMENTED,
    ↪augment_dog_car_squirrel, num_augmentations=2)
    apply_augmentations_for_class('car', CLIPS, AUGMENTED,
    ↪augment_dog_car_squirrel, num_augmentations=2)
    apply_augmentations_for_class('squirrel', CLIPS, AUGMENTED,
    ↪augment_dog_car_squirrel, num_augmentations=2)
    apply_augmentations_for_class('bird', CLIPS, AUGMENTED, augment_bird,
    ↪num_augmentations=4)
```

```
Augmenting dog: 100%|      | 2641/2641 [00:41<00:00, 63.89it/s]
Augmenting car: 100%|      | 2613/2613 [02:52<00:00, 15.16it/s]
Augmenting squirrel: 100%|    | 2479/2479 [00:38<00:00, 64.19it/s]
Augmenting bird: 100%|      | 1579/1579 [00:27<00:00, 56.97it/s]
```

Encapsulamos únicamente las clases que tienen una muestra representativa

```
[67]: allowed_classes = [
        'opossum', 'rabbit', 'coyote', 'cat', 'squirrel', 'raccoon',
        'dog', 'bobcat', 'car', 'bird'
        # NOTA: ahora estamos excluyendo fox, badger, deer, rodent, skunk
    ]
```

```
[70]: import shutil
import random
```

```
[71]: split_ratio = 0.25 # 25% valid, 75% train
```

```
[74]: for class_name in allowed_classes:
        src_dir = os.path.join(CLIPS, class_name)
        train_dir = os.path.join(DATALOADER, 'train', class_name)
        valid_dir = os.path.join(DATALOADER, 'valid', class_name)

        os.makedirs(train_dir, exist_ok=True)
        os.makedirs(valid_dir, exist_ok=True)
```



```

if not os.path.exists(src_dir):
    print(f"La carpeta {src_dir} no existe, saltando...")
    continue

files = [f for f in os.listdir(src_dir) if os.path.isfile(os.path.
↪join(src_dir, f))]
random.shuffle(files)

n_valid = int(len(files) * split_ratio)
valid_files = files[:n_valid]
train_files = files[n_valid:]

for f in tqdm(train_files, desc=f"{class_name} - train", leave=True):
    src = os.path.join(src_dir, f)
    dst = os.path.join(train_dir, f)
    shutil.copy2(src, dst)

for f in tqdm(valid_files, desc=f"{class_name} - valid", leave=True):
    src = os.path.join(src_dir, f)
    dst = os.path.join(valid_dir, f)
    shutil.copy2(src, dst)

print("Distribución train/valid completada.")

```

```

opossum - train: 100%|      | 8978/8978 [02:05<00:00, 71.40it/s]
opossum - valid: 100%|     | 2992/2992 [00:42<00:00, 71.02it/s]
rabbit - train: 100%|      | 3747/3747 [00:50<00:00, 74.62it/s]
rabbit - valid: 100%|     | 1248/1248 [00:16<00:00, 75.02it/s]
coyote - train: 100%|      | 3525/3525 [00:50<00:00, 70.03it/s]
coyote - valid: 100%|     | 1175/1175 [00:17<00:00, 69.01it/s]
cat - train: 100%|        | 3228/3228 [00:44<00:00, 73.15it/s]
cat - valid: 100%|       | 1075/1075 [00:15<00:00, 69.50it/s]
squirrel - train: 100%|   | 1860/1860 [00:23<00:00, 78.91it/s]
squirrel - valid: 100%|  | 619/619 [00:07<00:00, 79.82it/s]
raccoon - train: 100%|   | 5307/5307 [01:09<00:00, 76.61it/s]
raccoon - valid: 100%|  | 1769/1769 [00:22<00:00, 78.29it/s]
dog - train: 100%|       | 1981/1981 [00:25<00:00, 78.38it/s]
dog - valid: 100%|      | 660/660 [00:08<00:00, 79.45it/s]
bobcat - train: 100%|    | 3240/3240 [00:39<00:00, 82.15it/s]
bobcat - valid: 100%|   | 1080/1080 [00:13<00:00, 81.40it/s]
car - train: 100%|      | 1960/1960 [00:30<00:00, 63.53it/s]
car - valid: 100%|     | 653/653 [00:10<00:00, 61.99it/s]
bird - train: 100%|     | 1185/1185 [00:14<00:00, 83.28it/s]
bird - valid: 100%|    | 394/394 [00:04<00:00, 92.12it/s]

```

Distribución train/valid completada.


```
[77]: from vision.config import TRAIN, VALID
```

```
[78]: counts = count_files_per_class(TRAIN)
counts
```

```
[78]: {'opossum': 8978,
      'raccoon': 5307,
      'rabbit': 3747,
      'coyote': 3525,
      'bobcat': 3240,
      'cat': 3228,
      'dog': 1981,
      'car': 1960,
      'squirrel': 1860,
      'bird': 1185}
```

```
[79]: counts = count_files_per_class(VALID)
counts
```

```
[79]: {'opossum': 2992,
      'raccoon': 1769,
      'rabbit': 1248,
      'coyote': 1175,
      'bobcat': 1080,
      'cat': 1075,
      'dog': 660,
      'car': 653,
      'squirrel': 619,
      'bird': 394}
```

```
[ ]: from vision.config import AUGMENTED
```

```
[82]: counts = count_files_per_class(AUGMENTED)
counts
```

```
[82]: {'bird': 6316, 'dog': 5282, 'car': 5226, 'squirrel': 4958}
```

```
[80]: for cls in allowed_classes:
      augmented_dir = os.path.join(AUGMENTED, cls)
      train_dir = os.path.join(TRAIN, cls)
      os.makedirs(train_dir, exist_ok=True)

      if os.path.exists(augmented_dir):
          for fname in tqdm(os.listdir(augmented_dir), desc=f"{cls} - augmented",
                              leave=True):
              src = os.path.join(augmented_dir, fname)
              dst = os.path.join(train_dir, fname)
              shutil.copy2(src, dst)
```

```
squirrel - augmented: 100%|      | 4958/4958 [00:43<00:00, 113.12it/s]
dog - augmented: 100%|      | 5282/5282 [00:45<00:00, 116.22it/s]
car - augmented: 100%|      | 5226/5226 [01:20<00:00, 65.09it/s]
bird - augmented: 100%|      | 6316/6316 [00:52<00:00, 119.78it/s]
```

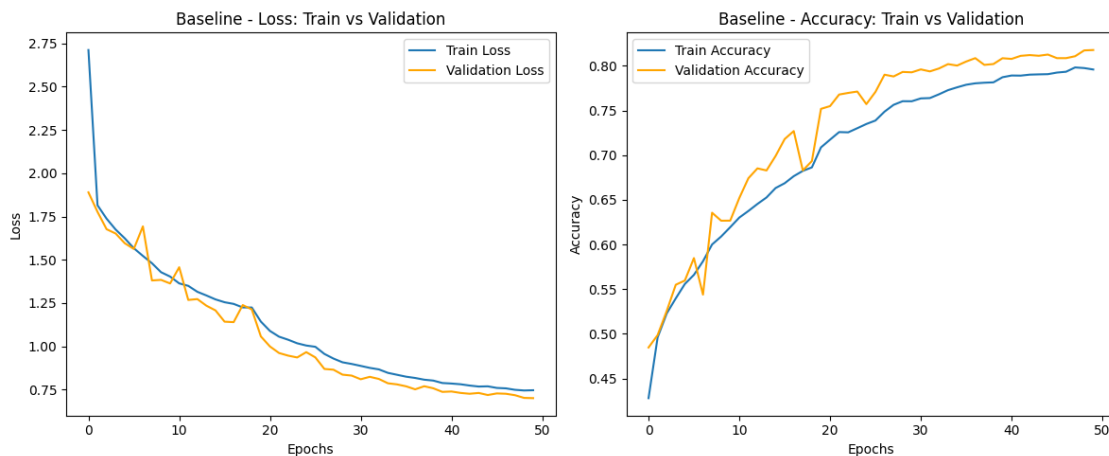
```
[81]: counts = count_files_per_class(TRAIN)
      counts
```

```
[81]: {'opossum': 8978,
      'bird': 7501,
      'dog': 7263,
      'car': 7186,
      'squirrel': 6818,
      'raccoon': 5307,
      'rabbit': 3747,
      'coyote': 3525,
      'bobcat': 3240,
      'cat': 3228}
```

```
[ ]: from vision.config import CHECKPOINTS_DIR
     from vision.plots import plot_training_history_from_dict
```

```
[2]: import json
     with open(CHECKPOINTS_DIR / "baseline_model_history.json", "r") as f:
         history_data = json.load(f)
```

```
[3]: plot_training_history_from_dict(history_data, title_prefix="Baseline - ")
```



modelo baseline

- **Loss:** Disminuye de forma estable en entrenamiento y validación, sin señales de sobreajuste.

- **Accuracy:** La precisión en validación supera a la de entrenamiento desde la época 10, lo que sugiere buena generalización.
- **Regularización:** El uso de Dropout y L2 fue efectivo para controlar el sobreentrenamiento.
- **Callbacks:** `EarlyStopping` y `ReduceLROnPlateau` funcionaron correctamente, estabilizando el entrenamiento.
- **Arquitectura:** La red CNN secuencial es adecuada como línea base para clasificación multiclase de fauna.

Conclusión: El modelo presentó un desempeño estable y generalizable, siendo un punto de partida sólido para futuras mejoras.

0.1.4 Desventajas del modelo actual

- Parte de cero → aprende bordes, texturas y patrones sin conocimiento previo.
 - Capacidad limitada → no captura la complejidad de fauna salvaje (variaciones de especies, fondos, luz).
 - Requiere muchos más datos y epochs para alcanzar niveles competitivos.
-

Recomendación - Cambiar a un modelo preentrenado (EfficientNetB0 o ResNet50) con fine-tuning.
 - Incorporar `class_weights` para compensar el desbalance entre clases. - Considerar optimizaciones como learning rate scheduler.

[]:

[]:

[]:

[]:

[]:

[]: